Reinforcement Learning

- Q - Learning
  - Off-Policy Temporal-Difference Control
    - differentiate behavior-policy from learning-policy
    - SARSA (on policy)
      - $<s, a, r, s', a' \leftarrow \pi(s)>$ => Learning
    - Q-Learning
      - $<s, a \leftarrow \pi(s), r, s'>$ => Learning
  - Update rule

$$Q(S_t, A_t)_{new} = Q(S_t, A_t)_{old} + \alpha[R_{t+1} + \gamma max_{a \in A} Q(S_{t+1}, a') - Q(S_t, A_t)_{old}]$$

Reinforcement Learning

- Function Approximation

  - Why function approximation?

    - Problem with large state spaces
      - Large memory for large table task
      - data should be accurate

    - Generalization
      - to generalize from previous encounters with different states that are in some sense similar to the current ones

  - Generalization => function approximation

    - to generalize desired functions (e.g value function, q function etc.)
    - utilize *supervised learning*